

data_exploration_features

January 26, 2017

0.1 Loading data/libs

```
In [1]: import pandas as pd
import numpy as np
import calendar

from bokeh.charts import output_notebook, Scatter, Bar, show, output_file, Line, BoxPlot
from bokeh.plotting import figure
from bokeh.layouts import row, column, gridplot

from ipywidgets import interactive
from IPython.display import display
from IPython.utils.py3compat import annotate

from graph import build_hist, to_relative_time
from fft import to_fft

output_notebook()
```

```
In [2]: INPUT="data/device_failure.csv"
dataset = pd.read_csv(INPUT, index_col=[0,1], parse_dates=[0])
```

features

Per features:

Statistical distribution: - Distribution - Distribution over failures - Distribution over devices - Distribution over failing devices

0.1.1 Temporal distribution

- Average value over time
- Average value over time for failing devices
- Value before failure

0.1.2 Frequency distribution

- DFT / device
- DFT / failing device

```
In [3]: @annotate(attribute=list(dataset.columns[1:]))
def pick_attribute(attribute):
    return "current attribute=%s" % s
s = interactive(pick_attribute)
display(s)

'current attribute=<ipywidgets.widgets.widget_box.Box object at 0x7fa41134de50>'
```

0.2 building data objects

```
In [4]: attribute = s.children[0].value
feature_dset = dataset[[attribute,"failure"]]
failing_points = feature_dset[feature_dset["failure">>0]
```

```
In [5]: def failure_date(failure):
    data = feature_dset.ix[failure.index]
    dates = data[data["failure">>0]
    if not dates.empty:
        return dates.iloc[0].name[0]
    else:
        return None
```

```
devices = feature_dset.groupby(level=1).agg(
    {
        "failure":{
            "failure":np.sum,
            "failure_date":failure_date},
        attribute : {
            "min_att":np.min,
            "max_att":np.max,
            "mean_att":np.mean,
            "std_att":np.std
        }
    })
devices.columns = devices.columns.droplevel()

failing_devices = devices[devices["failure">>0]
working_devices = devices[devices["failure"]==0]
```

```
working_devices_t = pd.DataFrame({attribute:feature_dset[attribute].unstack().filter(items=working_devices_t.columns)}
failing_devices_t = pd.DataFrame({attribute:feature_dset[attribute].unstack().filter(items=failing_devices_t.columns)

```

```
In [6]: p0 = build_hist(failing_points,attribute,u"%s for failing points" % attribute,color="green")
p1 = build_hist(feature_dset,attribute)

h = row(p0,p1)
show(h)
```

```
In [7]: def per_col(car):
        p0 = build_hist(failing_devices,car,label="%s on failing devices"% car)
        p1 = build_hist(working_devices,car,label="%s on working devices"% car)
        return [p0,p1]

        show(gridplot([ per_col(c) for c in set(c for c in devices.columns if "att" in c ) ]))
```

attribute 5: some failing points might be controlled by too much variation in this attribute

0.3 Average Value over Time.

```
In [8]: def dev_std(device_df):
        return device_df.rolling(window=20,center=False).std()

        def roll_std(df):
            return df[attribute].groupby(level="device").transform(dev_std)

        failing_devices_t["rolling_std"] = roll_std(failing_devices_t)
        working_devices_t["rolling_std"] = roll_std(working_devices_t)

        time = pd.DataFrame({
            attribute : feature_dset[attribute].groupby(level=0).mean(),
            "%s for failing devices" %attribute : failing_devices_t[attribute].groupby(level=1).
            "%s for working devices" %attribute : working_devices_t[attribute].groupby(level=1).
            "rolling std(%s) for failing devices" % attribute : failing_devices_t["rolling_std"]
            "rolling std(%s) for working devices" % attribute : working_devices_t["rolling_std"]
        })
        show(Line(time))
```

attribute1 seems to be completely different at the end of the period:

- Hyp 1: amplitude gets higher when the device start failing => disproved
- Hyp 2: amplitude is always higher for more fragile devices => disproved
- Hyp 2: wider amplitude for SOME signals
- Hyp 3: failing devices somehow synchronize, resonance effect (unlikely. plus what would it mean?)
- Hyp 4 : too few devices to average out see graph "n_devices" [here][1] [1]: data_exploration.ipynb

0.3.1 Watching samples aligned on the failure time

```
In [9]: rel_time_threshold = -100
        n_samples = 50

        # relative time is relative to the failure date for negatives
        fail_end_dates= devices["failure_date"].dropna().to_dict()
        fail_relative_time = to_relative_time(failing_devices_t,fail_end_dates,rel_time_threshold)
```

```

# for working ones, we use the last value (beware, could lead to weird effects, if the a
work_end_dates= working_devices_t.reset_index(level="date")["date"].groupby(level=0).max
work_relative_time = to_relative_time(working_devices_t,work_end_dates,rel_time_threshol

fail_rel_sampled = fail_relative_time[attribute].unstack(level="dt_from_fail").sample(n=
work_rel_sampled = work_relative_time[attribute].unstack(level="dt_from_fail").sample(n=

show(row(
    Line(
        fail_rel_sampled.unstack(level="device"),
        width=450,
        height=400,
        title = "%s before failure for a sample failing devices" % attribute,
        legend=None),
    Line(
        work_rel_sampled.unstack(level="device"),
        width=450,
        height=400,
        title = "%s before end for a sample working devices" % attribute,
        legend=None)
))

```

attribute 4: Need to take into account the Dv aver time, in addition to the dt

attribute 6 : we can see two classes of population: the ones with increasing attr6, and the ones without

0.3.2 Sampled devices in actual time

```

In [10]: n_samples=20
sampled_working_devices = working_devices_t[attribute].unstack(level="date").sample(n=n
sampled_failing_devices = failing_devices_t[attribute].unstack(level="date").sample(n=n

10 = Line(
    sampled_failing_devices.unstack(level="device"),
    width=450,
    height=400,
    title='%s for sampled failing devices' % attribute,
    legend=None
)
11 = Line(
    sampled_working_devices.unstack(level="device"),
    width=450,
    height=400,
    title = '%s for sampled working devices' % attribute,
    legend=None
)

show(row(10,11))

```

0.3.3 FFT

```
In [11]: fft_df = feature_dset[[attribute]].copy()
        fft_per_device = fft_df[attribute].groupby(level="device", sort=True).transform(to_fft)
        fft_df["df"] = fft_per_device

        fft_plot = fft_df.groupby(level="device").apply(lambda x: x.reset_index(drop=True))["df"]

        fft_working_devices = fft_plot.unstack(level=0).filter(items=working_devices.index).stack()
        fft_failing_devices = fft_plot.unstack(level=0).filter(items=failing_devices.index).stack()

        n_samples = 100
        to_plot_working = fft_working_devices.unstack(level=1).sample(n=10).unstack().dropna()
        to_plot_failing = fft_failing_devices.unstack(level=0).sample(n=10).unstack().dropna()
        show(row(
            Line(to_plot_failing.unstack("device"),
                width=450,
                height=400,
                title = "dft of %s for sampled failing devices" % attribute,
                legend=None),
            Line(to_plot_working.unstack("device"),
                width=450,
                height=400,
                title = "dft of %s for sampled working device" % attribute,
                legend=None),
        ))
```

attribute 6 : kampai !! attribute 9 : good

```
In [12]: if attribute == u"attribute3":
        dd = feature_dset
        dd[dd[attribute]>0]
        dd = dd.swaplevel().sort_index()
        grouped= dd[dd[attribute]>0].groupby(level="device").agg(lambda x : len(np.unique(x)))
        weird_devices = set(grouped[grouped>1].index)
        print weird_devices
        weird_values = dd.unstack(level="date").filter(items=weird_devices, axis="index").stack()
        print weird_values["failure"].value_counts()
        print weird_values.groupby(level="device").apply(lambda df: df[attribute].value_counts())

In [13]: # attribute 5: testing if variations are not wider for some failures..
        print (failing_devices["max_att"] -failing_devices["min_att"]).value_counts()
        print (working_devices["max_att"] -working_devices["min_att"]).value_counts()

        # answer: Nope. ...
```

0	62
8	2
2336	1

9264	1
160	1
328	1
10288	1
168	1
32	1
88	1
64712	1
368	1
448	1
648	1
62488	1
16	1
63504	1
2576	1
1208	1
520	1
1656	1
21816	1
920	1
48	1
80	1
2168	1
61088	1
240	1
10336	1
27856	1
4360	1
2792	1
2144	1
136	1
2496	1
10200	1
304	1
2408	1
21200	1
976	1
64	1
10440	1
456	1
960	1
dtype: int64	
0	1003
8	8
16	5
296	3
40	3
96	2

```

144      2
464      2
32       2
24       2
112      1
120      1
416      1
55464    1
176      1
200      1
208      1
8544     1
49192    1
128      1
3976     1
1872     1
912      1
1752     1
3696     1
1384     1
64792    1
1240     1
64584    1
1000     1
2864     1
552      1
51976    1
4832     1
6856     1
2688     1
632      1
616      1
2648     1
424      1
dtype: int64

```

```
In [14]: ## Isolate a validation dataset
```

```
# easier without index
```

```
all_devices = dataset.reset_index()[["device", "failure"]]
```

```
#Keep 5% of each class for validating
```

```
validation_neg = set(all_devices[all_devices["failure"]==1].sample(frac=0.05, random_st
```

```
validation_pos = set(all_devices[all_devices["failure"]==0].sample(frac=0.05, random_st
```

```
validation_devices = validation_neg.union(validation_pos)
```

```

# build the two sets
per_device = dataset.reset_index()
per_device["k"] = per_device["device"].map(dict( (k,True) for k in validation_devices))
per_device["k"].fillna(False,inplace=True)
g = per_device.groupby(by="k")
training_set = g.get_group(True).set_index(drop=True,keys=["date","device"])
validation_set = g.get_group(False).set_index(drop=True,keys=["date","device"])
del training_set["k"]
del validation_set["k"]

# save the two sets
print validation_set.shape
print training_set.shape
validation_set.to_csv("data/validation.csv")
training_set.to_csv("data/train.csv")

```

```

(2729, 10)
(121765, 10)

```

```
In [ ]:
```